



R U N T I M E M A I N T E N A N C E L A Y E R F O R E D G E A I



The right-sized AI model, on the right chip, at the right time

The Haiku-not-Opus moment for industrial AI

Kevin Omede · Politecnico di Torino × UIC Barcelona · Visiting Researcher, Beihang University

PROBLEM

\$2.3M per hour It's the cost of
automotive downtime when AI misses the alarm.

3–6 months

average delay before silent
degradation is detected

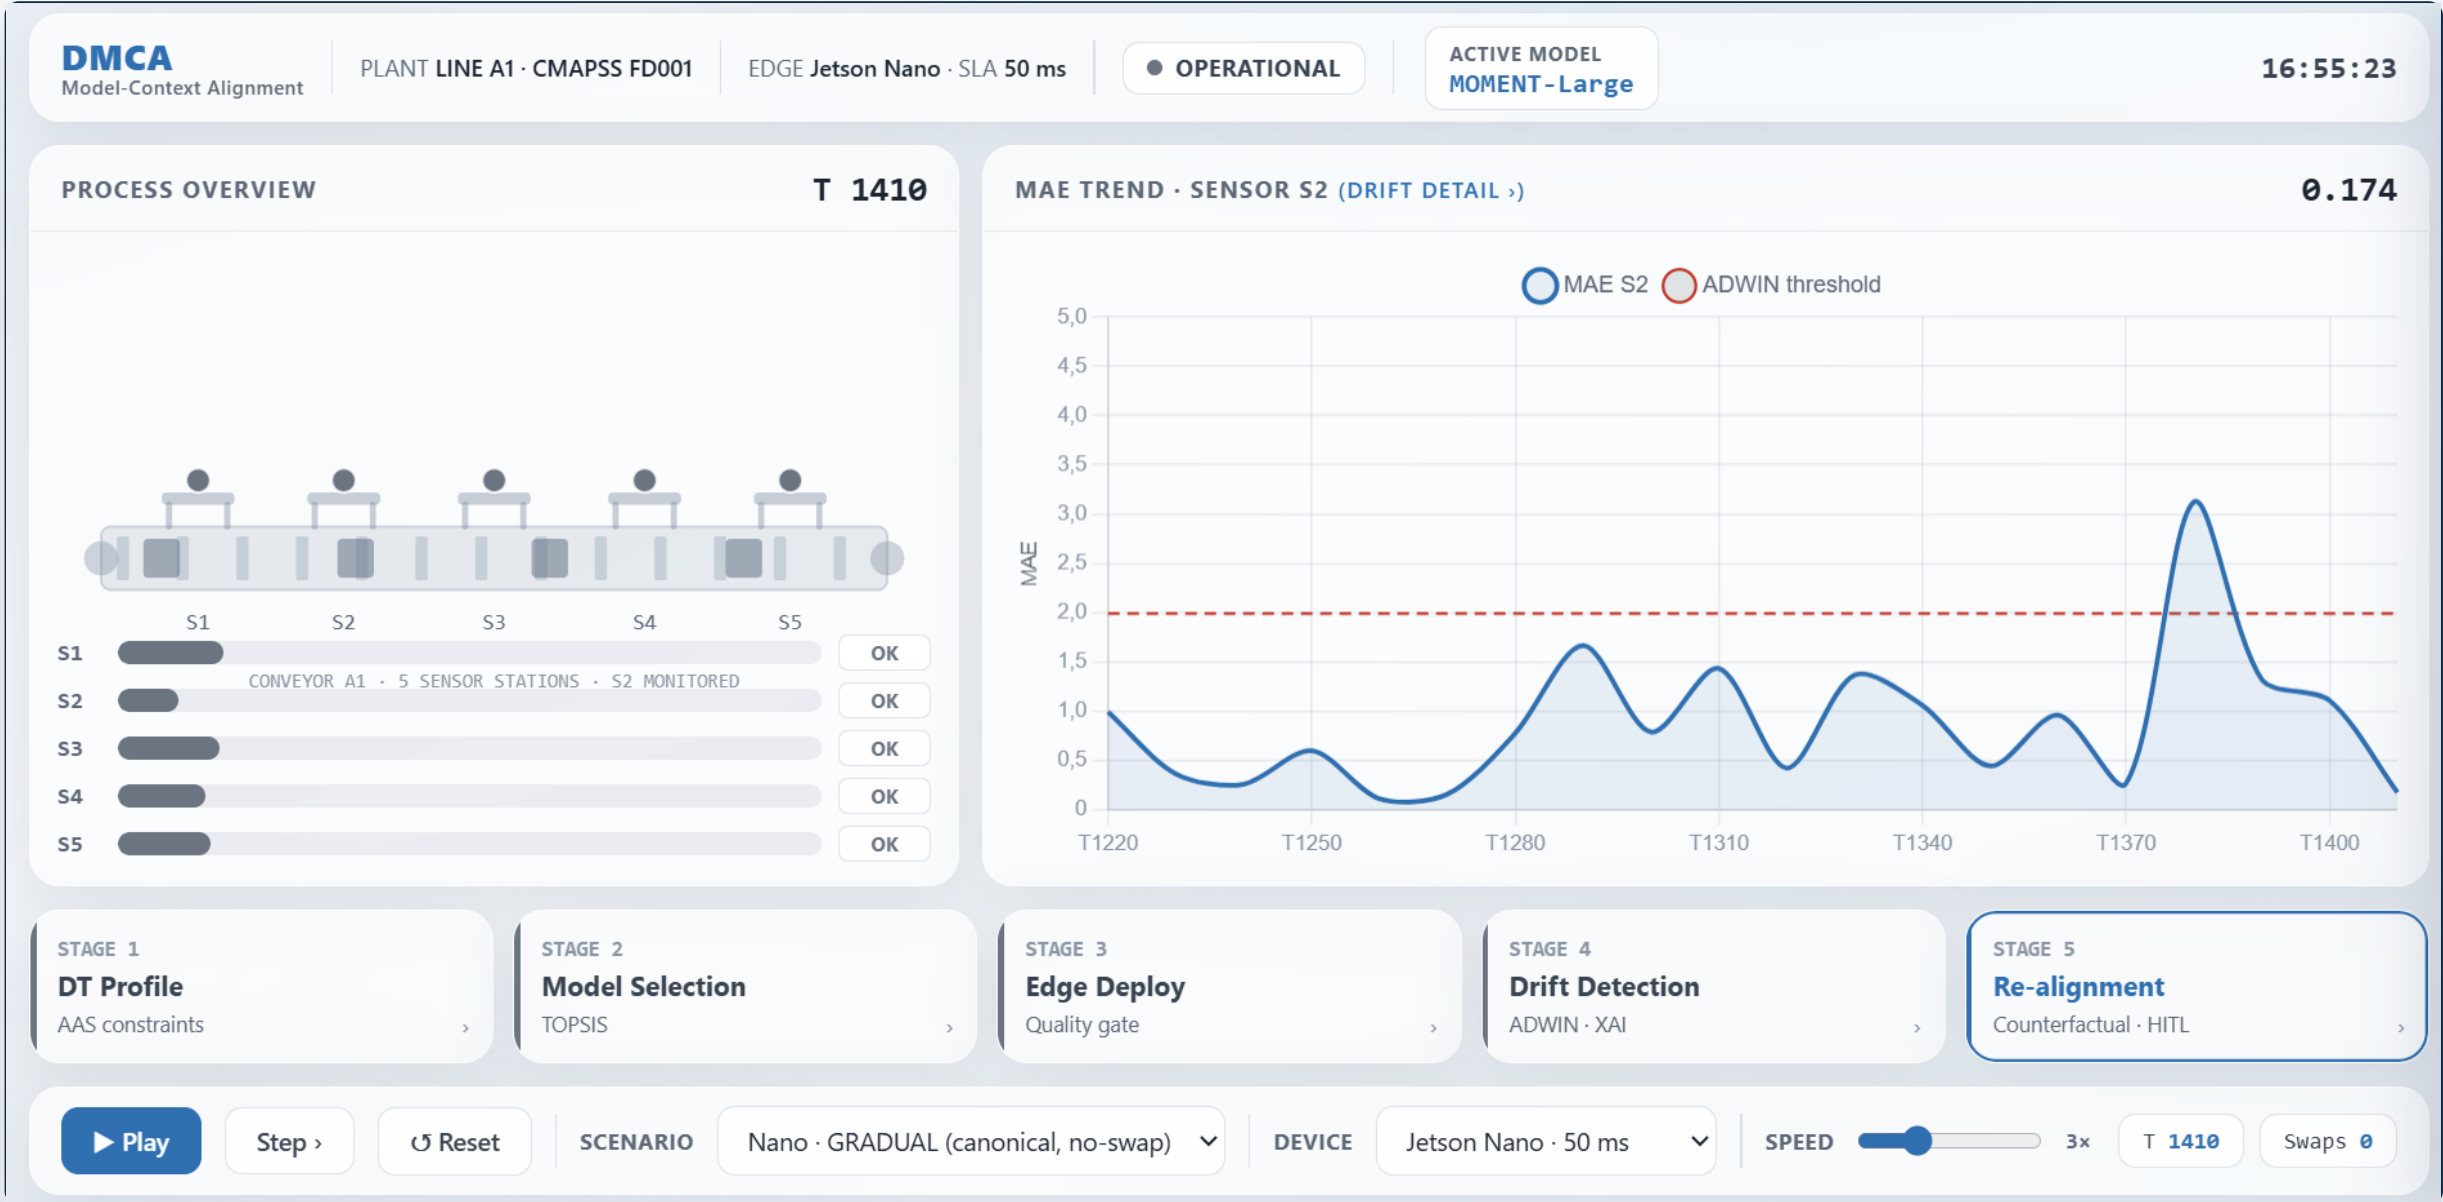
Invisible

the cloud never sees SLA-violating
requests, so missed alarms never
reach monitoring

Unmeasured

nobody checks the model actually
runs in time on the real chip

SOLUTION We test your models on real edge hardware — then keep them working



01 You send

Your model, a data sample, hardware spec, SLA budget.

02 We test

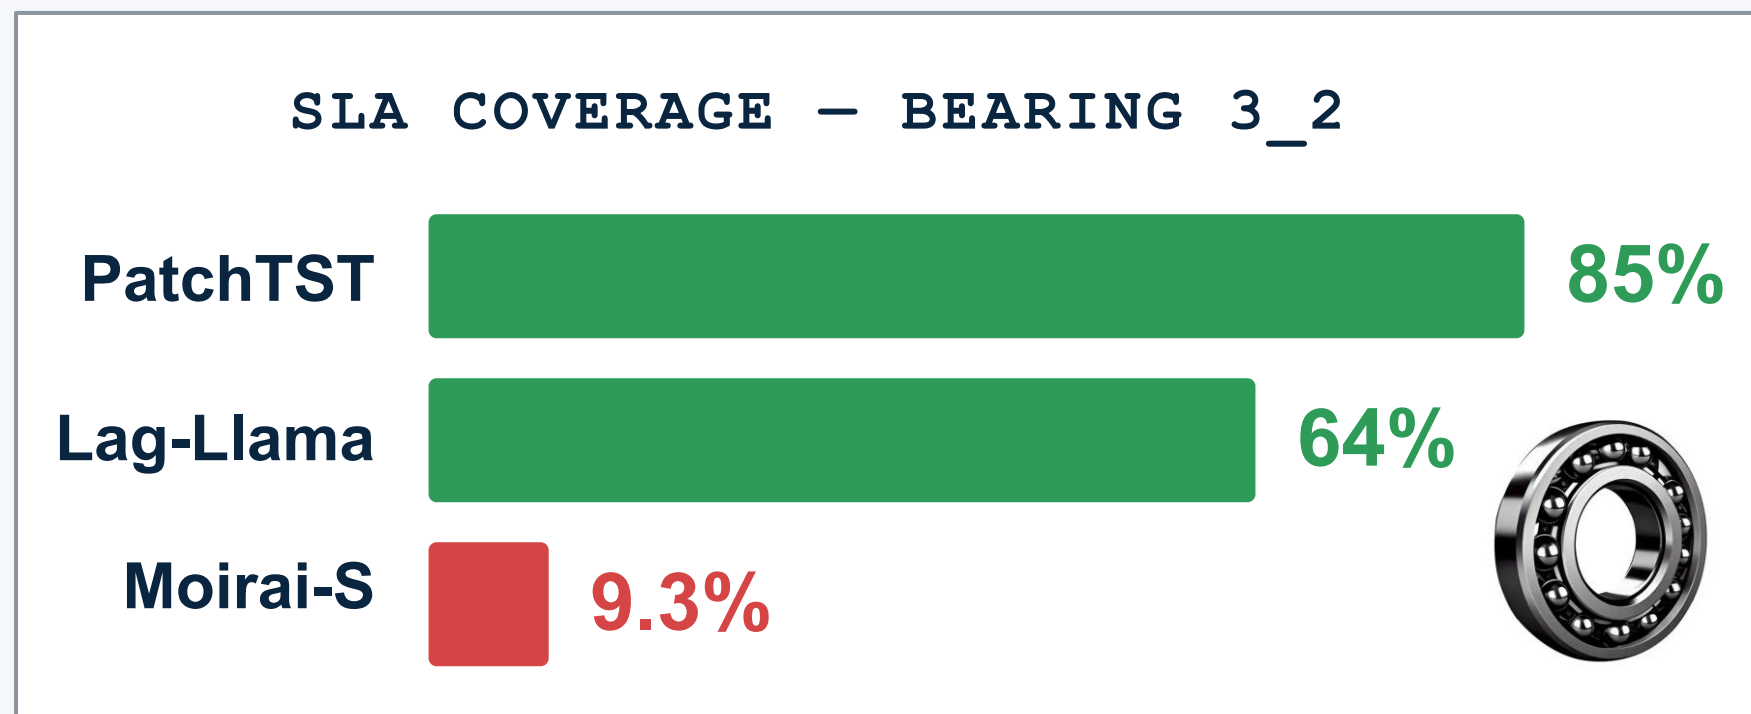
Real-hardware profiling, drift injection, coverage against your SLA and TOPSIS vs 8 foundation models.

03 You get

Full report telling you which models fit, which don't and how to manage them.

Everyone measures accuracy.

Nobody measures coverage.



10x gap in SLA coverage
same task, same hardware, same SLA

WHY NOW

Three converging shifts open a 24-month window

REGULATORY

3%

of global turnover.

EU AI Act fines hit August 2, 2026.

FOUNDATION MODEL EXPLOSION

2.4M+

models on HuggingFace.

Manual selection is now impossible.

CLOSED ECOSYSTEMS

LOCK-IN

Every chip vendor locks you in. No vendor has incentive to build the agnostic decision layer.

Regulation demands it. Complexity requires it. No one is building it.

BUSINESS MODEL

One wedge, one core, one platform

LAND

Audit

On-hardware validation.

First market entry.

EXPAND

Runtime

Continuous drift detection +
automatic swap.

Enterprise subscription.

DEEPEN

Optimization

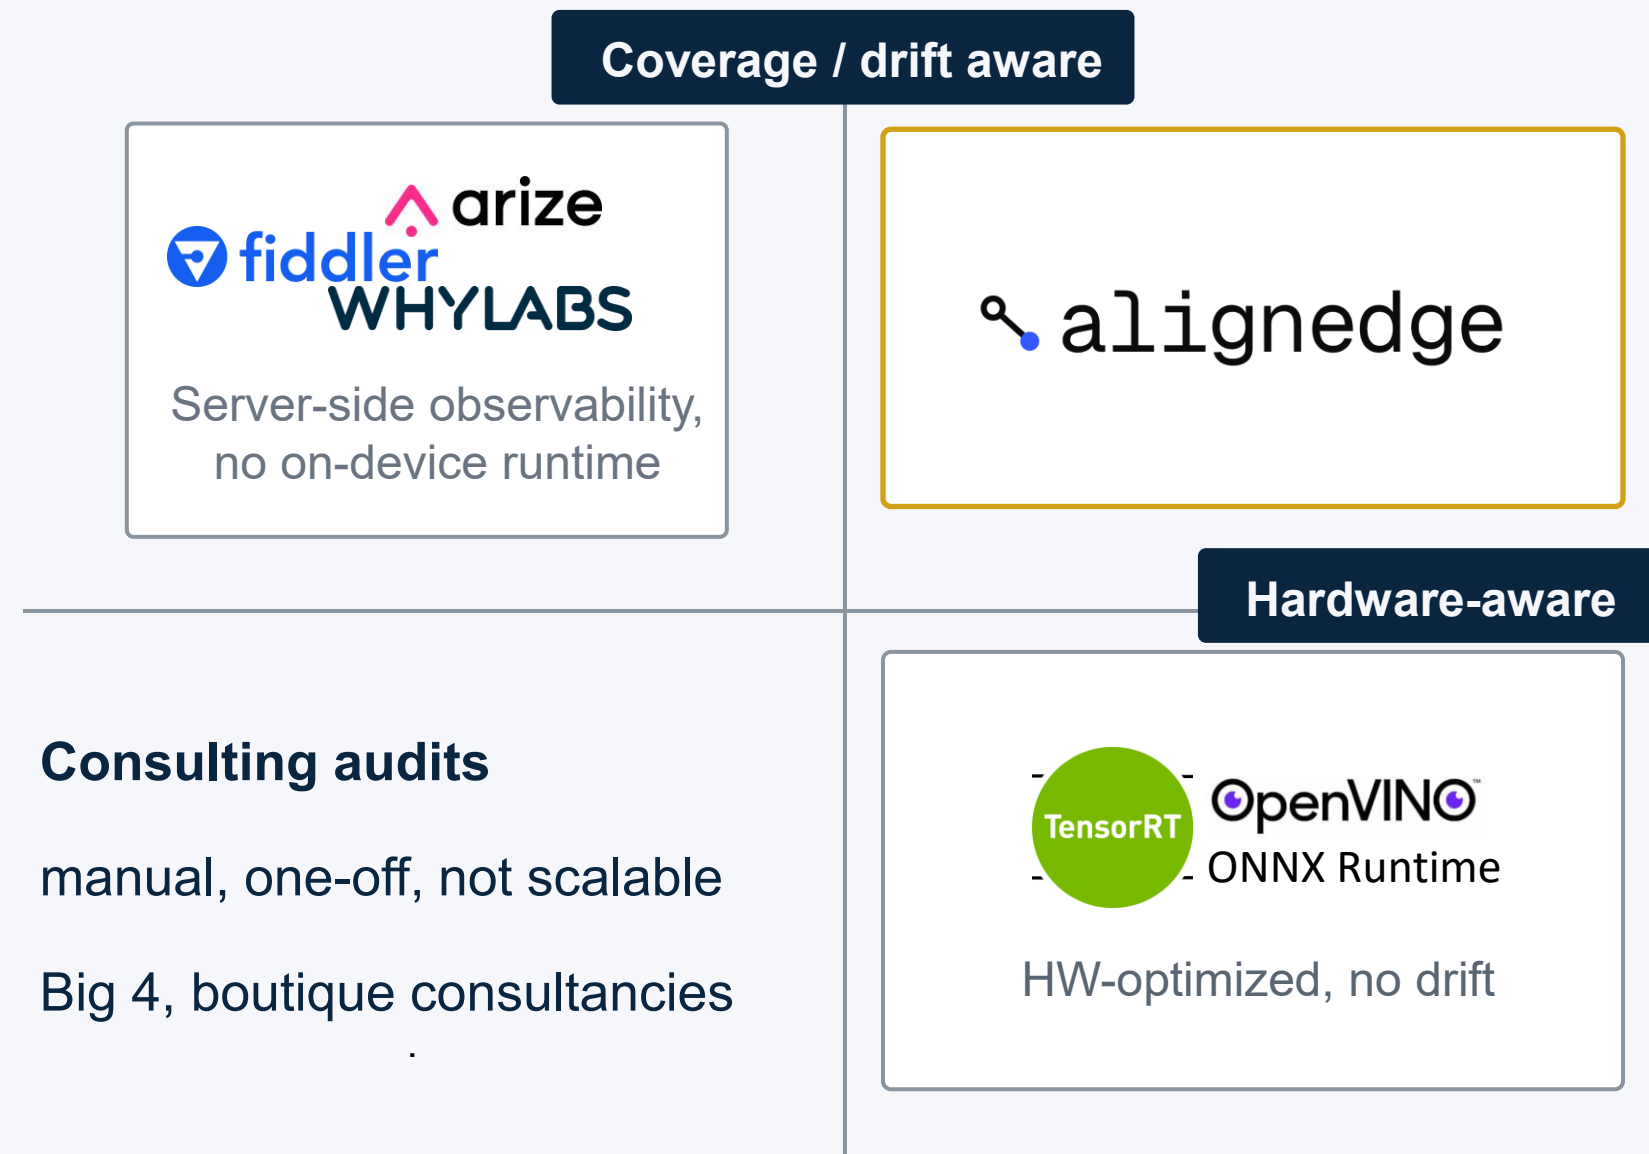
Fleet-wide cross-vendor
orchestration

Platform scale integration

Same customer, growing **footprint** over time.

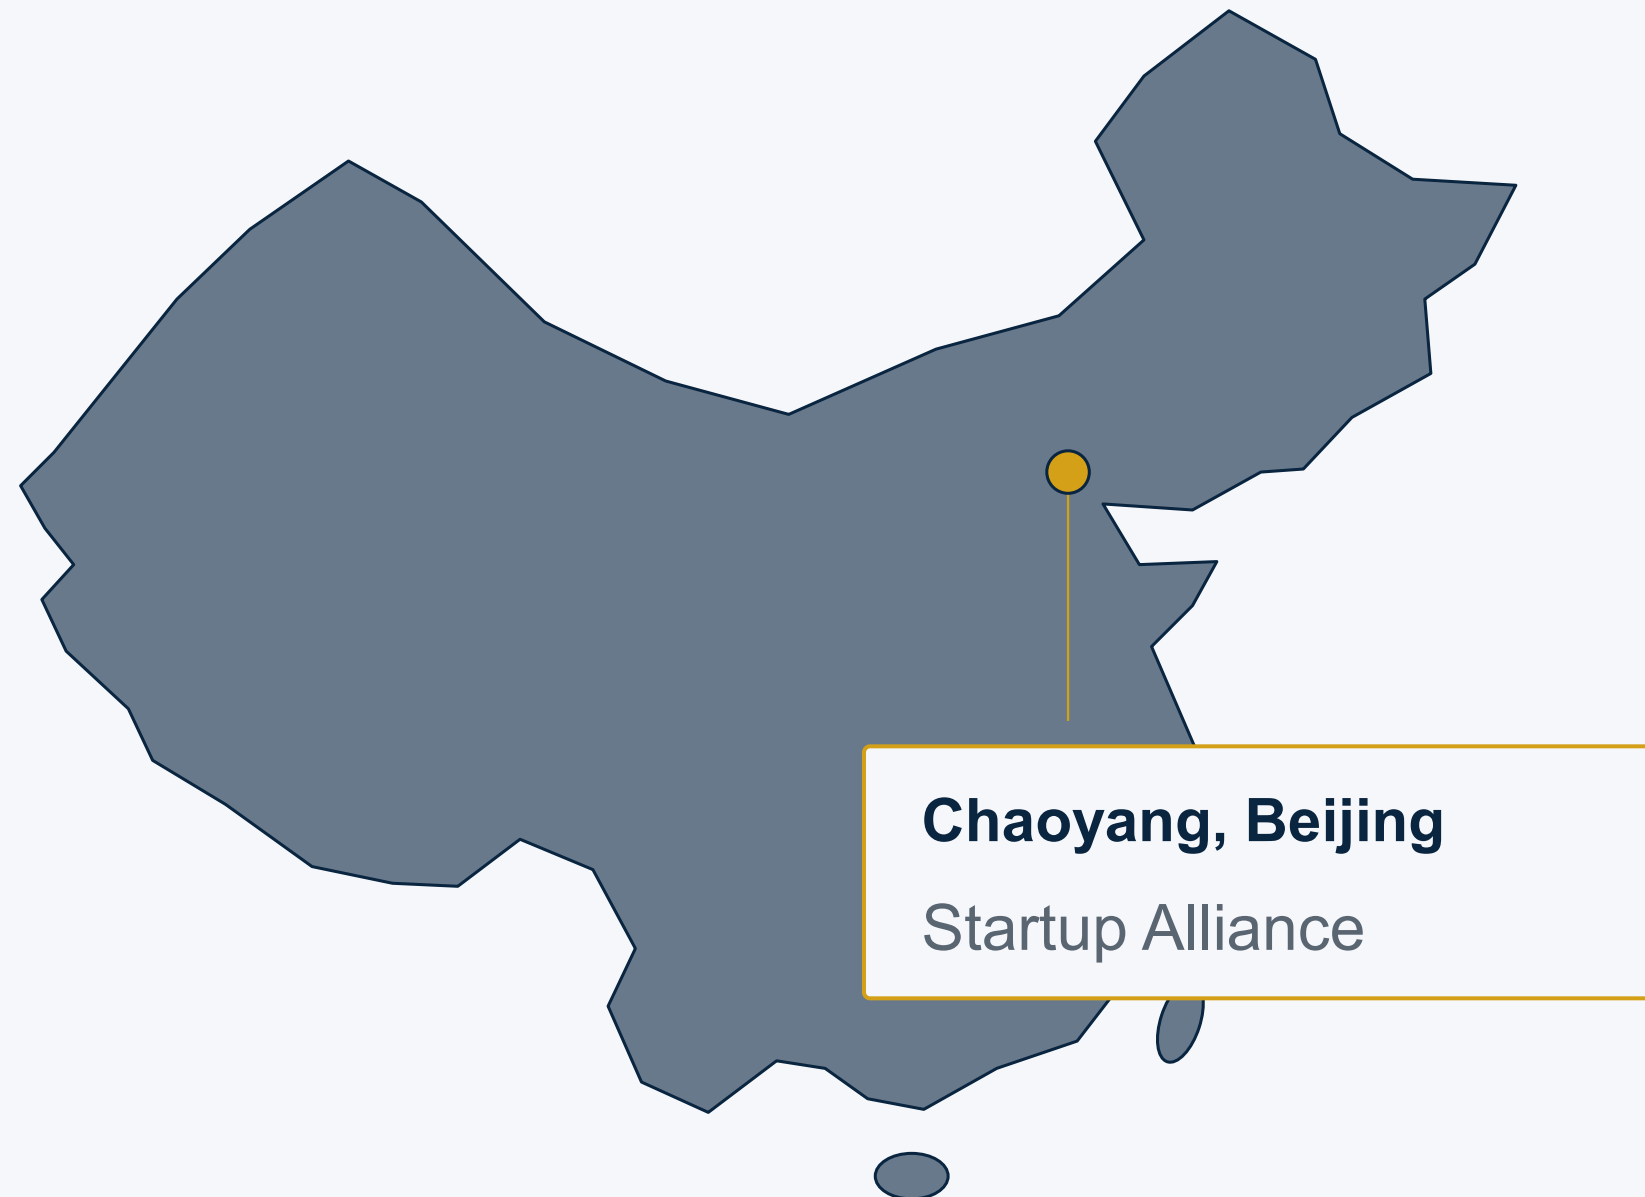
COMPETITION

Three reasons no one else can catch up



- 1 Productized audit measuring SLA coverage on your hardware, with your data.
- 2 Cross-vendor by design. No chip vendor will cannibalize their own lock-in.
- 3 A dataset that compounds — every audit adds one customer × model × hardware × drift × outcome record no competitor can replicate

CHINA · 中国



75%

of the world's EVs are produced in China

\$318B

global automotive electronics
market (2026)

43.6%

Asia Pacific share

The largest edge AI deployment testbed on Earth.

TEAM + TRACTION

One founder, three institutions, validated technology



Politecnico di Torino

UIC
barcelona

Kevin Omede

FOUNDER

MSc, Politecnico di Torino ×
Universitat Internacional de Catalunya
Visiting Researcher,
Beihang University

- 92% drift classifier accuracy, held-out
- 0% false alarms
- \$0.30 canonical cost — 8 models on real GPU
- 10× coverage gap — PHM 2012 + IMS Center

Every industrial AI model, always in its optimal deployment state

AI-on-edge becomes a maintained asset, not a one-shot deployment. By 2030: millions of models continuously maintained. Hundreds of billions of dollars in industrial downtime prevented.

The Haiku-not-Opus moment for industrial AI.